Paper List

whoisKIWIIZZZ October 24, 2026

Abstract

整理一系列文章为了后续能快速的找到内容.

1 World Model

1.1 JEPAs

Training object is (likely) representations, and use CEM to MPC.

- I-JEPA,V-JEPA [1, 2, 3, 4, 5]
- leWM-based [6, 7, 8],
- DINO-WM [9]
- [10], Hierarchical planning

1.2 latent representations

I think that JEPA also helps for building a 'good' representations of model prediction. Here are works that focusing on training the encoder instead of just next-step prediction.

- [11], smoothness for representations
- [12], project DINO token
- *[13], planning in latent*

1.3 WAMs

- [14], predict future
- [15], use optical-flow. *TODO:add lfy's work*
- [16], just denoising w/o predict future

1.4 Others

- [17],Point WMs
- [18], dataset for other goals
- [19], testing on Video-gen
- [12], use ODE as training objective
- [20], good pipeline on WM
- [21, 22], ARC-AGI (idk where to put it)
- [23], Probe to find whether WM can find the original laws(e.g.,gravity)

2 Loop

- [24],LoopViT
- [25], use DP to find training-free loops in ViTs
- [26], refine in Diffusion
- [27], LoopDiT
- [28], LoopWM

3 NCA

- [29], Summary
- [30], Diffusion as NCA
- [31, 32], 2D reasoning w/ NCA
- Yitao Xu[1]'s prior work:
  - Growing NCA [33], the original NCA
  - DyNCA [34], dynamic texture synthesis by motion supervision
  - MeshNCA [35], texture synthesis on meshes
  - [36], NCA + appearance decoder, from cells to high-res pixels
  - NoiseNCA [37], noisy seed for spatio-temporal continuity
  - [38], Neural Particle Automata

4 Theory

- [39],question ID/OOD acc using **gradient flow**. Check if u wish to see analyses of subspace
- [40],great work of "capacity" of models[2]
- several works of JEPAs.
  - [7],Gaussian world, how SIGReg work ?
  - [41], LpWM , sprase-space analysis
  - [11], smoothness for representations
  - [42], decode state of WMs to identify
- [43], random matrixs

[1] thx for his wonderful intuition!
[2] here to say, inductive bias is needed for models to learn

References

[1] Mido Assran, Adrien Bardes, David Fan, Quentin Garrido, Russell Howes, Mojtaba, Komeili, Matthew Muckley, Ammar Rizvi, Claire Roberts, Koustuv Sinha, Artem Zholus, Sergio Arnaud, Abha Gejji, Ada Martin, Francois Robert Hogan, Daniel Dugas, Piotr Bojanowski, Vasil Khalidov, Patrick Labatut, Francisco Massa, Marc Szafraniec, Kapil Krishnakumar, Yong Li, Xiaodong Ma, Sarath Chandar, Franziska Meier, Yann LeCun, Michael Rabbat, and Nicolas Ballas. V-JEPA 2: Self-Supervised Video Models Enable Understanding, Prediction and Planning. *arXiv preprint arXiv:2506.09985*, 2025.

[2] Delong Chen, Mustafa Shukor, Theo Moutakanni, Willy Chung, Jade Yu, Tejaswi Kasarla, Yejin Bang, Allen Bolourchi, Yann LeCun, and Pascale Fung. VL-JEPA: Joint Embedding Predictive Architecture for Vision-language. *arXiv preprint arXiv:2512.10942*, 2025.

[3] Adrien Bardes, Quentin Garrido, Jean Ponce, Xinlei Chen, Michael Rabbat, Yann LeCun, Mahmoud Assran, and Nicolas Ballas. Revisiting Feature Prediction for Learning Visual Representations from Video. *arXiv preprint arXiv:2404.08471*, 2024.

[4] Mahmoud Assran, Quentin Duval, Ishan Misra, Piotr Bojanowski, Pascal Vincent, Michael Rabbat, Yann LeCun, and Nicolas Ballas. Self-Supervised Learning from Images with a Joint-Embedding Predictive Architecture. *arXiv preprint arXiv:2301.08243*, 2023.

[5] Lorenzo Mur-Labadia, Matthew Muckley, Amir Bar, Mido Assran, Koustuv Sinha, Mike Rabbat, Yann LeCun, Nicolas Ballas, and Adrien Bardes. V-JEPA 2.1: Unlocking Dense Features in Video Self-Supervised Learning. *arXiv preprint arXiv:2603.14482*, 2026.

[6] Lucas Maes, Quentin Le Lidec, Damien Scieur, Yann LeCun, and Randall Balestriero. LeWorldModel: Stable End-to-End Joint-Embedding Predictive Architecture from Pixels. *arXiv preprint arXiv:2603.19312*, 2026.

[7] David Klindt, Yann LeCun, and Randall Balestriero. When Does LeJEPA Learn a World Model? *arXiv preprint arXiv:2605.26379*, 2026.

[8] Ying Wang, Oumayma Bounou, Yann LeCun, and Mengye Ren. AdaJEPA: An Adaptive Latent World Model. *arXiv preprint arXiv:2606.32026*, 2026.

[9] Gaoyue Zhou, Hengkai Pan, Yann LeCun, and Lerrel Pinto. DINO-WM: World Models on Pre-trained Visual Features enable Zero-shot Planning. *arXiv preprint arXiv:2411.04983*, 2024.

[10] Wancong Zhang, Basile Terver, Artem Zholus, Soham Chitnis, Harsh Sutaria, Mido Assran, Randall Balestriero, Amir Bar, Yann LeCun, and Nicolas Ballas. Hierarchical Planning with Latent World Models. *arXiv preprint arXiv:2604.03208*, 2026.

[11] Ying Wang, Oumayma Bounou, Gaoyue Zhou, Randall Balestriero, Tim G. J. Rudner, Yann LeCun, and Mengye Ren. Temporal Straightening for Latent Planning. *arXiv preprint arXiv:2603.12231*, 2026.

[12] Dongxiu Liu, Haoyi Niu, Peng Cheng, Yuan Gao, Xirui Kang, Sangli Teng, Koushil Sreenath, and Xianyuan Zhan. ODEWorld: A Continuous Predictive Architecture via Physical-Time Flow. *arXiv preprint arXiv:2607.27924*, 2026.

[13] Yijun Yuan, Weicheng Zheng, Weibang Wang, Minghui Qin, Chang Sun, Junhao Huang, Kenan Li, Anmin Liu, Yicheng Yao, and Hang Zhao. Representation World Model: Learning States, Transition and Executable Plans in Representation. *arXiv preprint arXiv:2609.29171*, 2026.

[14] Yuyang Zhang, Wenyao Zhang, Zekun Qi, He Zhang, Haitao Lin, Jingbo Zhang, Yao Mu, Xiaokang Yang, Wenjun Zeng, and Xin Jin. ImageWAM: Do World Action Models Really Need Video Generation, or Just Image Editing? *arXiv preprint arXiv:2606.19531*, 2026.

[15] Yixiang Chen, Peiyan Li, Yuan Xu, Qisen Ma, Jiabing Yang, Kai Wang, Jianhua Yang, Dong An, He Guan, Gaoteng Liu, Jianlou Si, Jun Huang, Jing Liu, Nianfeng Liu, Yan Huang, and Liang Wang. FlowWAM: Optical Flow as a Unified Action Representation for World Action Models. *arXiv preprint arXiv:2607.13017*, 2026.

[16] Zanyi Wang, Yuheng Lei, Dengyang Jiang, Ping Luo, Mengdi Wang, Zhixuan Liang, and Shilong Liu. Beyond Future Prediction: Denoising as Generative Adaptation for Robot Control. *arXiv preprint arXiv:2609.28339*, 2026.

[17] Tal Daniel, Carl Qi, Dan Haramati, Amir Zadeh, Chuan Li, Aviv Tamar, Deepak Pathak, and David Held. Latent Particle World Models: Self-supervised Object-centric Stochastic Dynamics Modeling. *arXiv preprint arXiv:2603.04553*, 2026.

[18] Heejeong Nam, Divyat Mahajan, Quentin Le Lidec, Randall Balestriero, and David Klindt. AnyGoal: On World Model Generalization Beyond Expert Rollouts. *arXiv preprint arXiv:2610.26844*, 2026.

[19] Weihang Guo, Xiaoyu Wu, Yifei Wang, Niloofar Mireshghallah, and Lydia E. Kavraki. Scaling Video Generation for Reasoning: At What Cost? *arXiv preprint arXiv:2609.36599*, 2026.

[20] Siqiao Huang, Partha Kaushik, Michael Chen, Hengkai Pan, Kaiwen Geng, Omar Chehab, Fernando Moreno-Pino, and Max Simchowitz. Nano World Models: A Minimalist Implementation of Future Video Prediction. *arXiv preprint arXiv:2605.23993*, 2026.

[21] Zhiguang Liu and Yi Shang. Reasoning is a Modality. *arXiv preprint arXiv:2601.13562*, 2026.

[22] Keya Hu, Ali Cy, Linlu Qiu, Xiaoman Delores Ding, Runqian Wang, Yeyin Eva Zhu, Jacob Andreas, and Kaiming He. ARC Is a Vision Problem! *arXiv preprint arXiv:2511.14761*, 2025.

[23] Keyon Vafa, Peter G. Chang, Ashesh Rambachan, and Sendhil Mullainathan. What Has a Foundation Model Found? Using Inductive Bias to Probe for World Models. *arXiv preprint arXiv:2507.06952*, 2025.

[24] Wen-Jie Shu, Xuerui Qiu, Rui-Jie Zhu, Harold Haodong Chen, Yexin Liu, and Harry Yang. LoopViT: Scaling Visual ARC with Looped Transformers. *arXiv preprint arXiv:2602.02156*, 2026.

[25] Mozes Jacobs, Thomas Fel, Richard Hakim, Alessandra Brondetta, Demba Ba, and T. Andy Keller. Block-Recurrent Dynamics in Vision Transformers. *arXiv preprint arXiv:2512.19941*, 2025.

[26] Sahil Goyal, Swayam Agrawal, Gautham Govind Anil, Prateek Jain, Sujoy Paul, and Aditya Kusupati. ELT: Elastic Looped Transformers for Visual Generation. *arXiv preprint arXiv:2604.09168*, 2026.

[27] Yong Xien Chng, Tianyi Chen, Wenwen Tong, Haiwen Diao, Zhongang Cai, Lei Yang, Ziwei Liu, Lewei Lu, Dahua Lin, and Gao Huang. Looped Diffusion Transformer. *arXiv preprint arXiv:2609.40305*, 2026.

[28] Hongyuan Adam Lu, Z. L. Victor Wei, Qun Zhang, Jinrui Zeng, Bowen Cao, Lingwei Meng, Mocheng Li, Zezhong Wang, Haonan Yin, Naifu Xue, Minyu Chen, Cenyuan Zhang, Zefan Zhang, Hao Wei, Jiawei Zhou, Haoran Xu, Hao Yang, Ronglai Zuo, Tongda Xu, Yonghao Li, Jian Chen, Hebin Wang, Zeyu Gao, Yang Li, Wei Zhao, Qimin Zhong, Siqi Liu, Yumeng Zhang, Leyan Cui, Zhangyu Wang, and Wai Lam. Looped World Models. *arXiv preprint arXiv:2606.18208*, 2026.

[29] Martin Spitznagel and Janis Keuper. A New Kind of Network? Review and Reference Implementation of Neural Cellular Automata. *arXiv preprint arXiv:2604.24990*, 2026.

[30] Mariia Drozdova, Aidan Sirbu, Pietro Miotti, Robert Obryk, Mayalen Etcheverry, Eyvind Niklasson, and Blake Richards. Diffusion as a Training Curriculum for Timestep-Free Iterative Reasoning. *arXiv preprint arXiv:2609.01449*, 2026.

[31] Martin Spitznagel and Janis Keuper. 2D Spatial Reasoning with Adaptive Neural Cellular Automata. *arXiv preprint arXiv:2610.08518*, 2026.

[32] Mayalen Etcheverry, Pietro Miotti, Aidan Sirbu, Konstantin Schürholt, Mariia Drozdova, Arna Ghosh, Blaise Agüera y Arcas, James Manyika, Blake Richards, and Eyvind Niklasson. Reasoning with Neural Cellular Automata. *arXiv preprint arXiv:2609.36126*, 2026.

[33] Alexander Mordvintsev, Ettore Randazzo, Eyvind Niklasson, and Michael Levin. Growing Neural Cellular Automata. *Distill* 5.2, 2020, e23.

[34] Ehsan Pajouheshgar, Yitao Xu, Tong Zhang, and Sabine Süsstrunk. "DyNCA: Real-time Dynamic Texture Synthesis Using Neural Cellular Automata". *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2023, pp. 20742–20751.

[35] Ehsan Pajouheshgar, Yitao Xu, Alexander Mordvintsev, Eyvind Niklasson, Tong Zhang, and Sabine Süsstrunk. Mesh Neural Cellular Automata. *ACM Transactions on Graphics (TOG)* 43.4, 2024, pp. 1–16.

[36] Ehsan Pajouheshgar, Yitao Xu, Ali Abbasi, Alexander Mordvintsev, Wenzel Jakob, and Sabine Süsstrunk. Neural Cellular Automata: From Cells to Pixels. *arXiv preprint arXiv:2506.22899*, 2025.

[37] Ehsan Pajouheshgar, Yitao Xu, and Sabine Süsstrunk. "NoiseNCA: Noisy Seed Improves Spatio-Temporal Continuity of Neural Cellular Automata". *Artificial Life Conference Proceedings 36*. Vol. 2024. MIT Press, 2024, p. 57.

[38] Hyunsoo Kim, Ehsan Pajouheshgar, Sabine Süsstrunk, Wenzel Jakob, and Jinah Park. Neural Particle Automata: Learning Self-Organizing Particle Dynamics. *arXiv preprint arXiv:2601.16096*, 2026.

[39] Ananya Kumar, Aditi Raghunathan, Robbie Jones, Tengyu Ma, and Percy Liang. Fine-Tuning can Distort Pretrained Features and Underperform Out-of-Distribution. *arXiv preprint arXiv:2202.10054*, 2022.

[40] Vassilis Papadopoulos, Jérémie Wenger, and Clément Hongler. Arrows of Time for Large Language Models. *arXiv preprint arXiv:2401.17505*, 2024.

[41] Yilun Kuang, Yash Dagade, Quentin Le Lidec, Lucas Maes, Randall Balestriero, and Yann LeCun. LpWM: A Case for Sparse Representations in World Models. *arXiv preprint arXiv:2608.22764*, 2026.

[42] Xiangteng Zhang, Yang Guan, Bo Zhang, Hongyang Li, Ya-Qin Zhang, and Shengbo Eben Li. On the Identifiability of Controlled World Models. *arXiv preprint arXiv:2607.22430*, 2026.

[43] Joel A. Tropp. User-friendly tail bounds for sums of random matrices. *arXiv preprint arXiv:1004.4389*, 2010.